

Matrix Algebra

Gaston Sanchez

November 7, 2016

Basic Vector and Matrix manipulations in R

If you have no or minimal experience in R working with matrices, please go over this section. If you are already familiar with this material, skip to the next section.

Consider the following vector `x`:

```
x <- 1:9
```

- Use the vector `x` as input of the function `matrix()` to create the following matrix:

	[,1]	[,2]	[,3]
[1,]	1	4	7
[2,]	2	5	8
[3,]	3	6	9

- Using `x` and `matrix()`, how would you generate a matrix like this:

	[,1]	[,2]	[,3]
[1,]	1	2	3
[2,]	4	5	6
[3,]	7	8	9

- Use `diag()` to create the following identity matrix \mathbf{I}_n of dimensions $n \times p = (5, 5)$:

	[,1]	[,2]	[,3]	[,4]	[,5]
[1,]	1	0	0	0	0
[2,]	0	1	0	0	0
[3,]	0	0	1	0	0
[4,]	0	0	0	1	0
[5,]	0	0	0	0	1

- Consider the following vectors `a1`, `a2`, `a3`:

```
a1 <- c(2, 3, 6, 7, 10)
a2 <- c(1.88, 2.05, 1.70, 1.60, 1.78)
a3 <- c(80, 90, 70, 50, 75)
```

- Column-bind the vectors `a1`, `a2`, `a3` to form this matrix:

	a1	a2	a3
1	2	1.88	80
2	3	2.05	90
3	6	1.70	70
4	7	1.60	50
5	10	1.78	75

- Now row-bind the vectors `a1`, `a2`, `a3` to form this matrix:

```

      1      2      3      4      5
a1  2.00  3.00  6.0  7.0 10.00
a2  1.88  2.05  1.7  1.6  1.78
a3 80.00 90.00 70.0 50.0 75.00

```

- Using matrix functions like transpose `t()` and product `%*%`, write a function `vnorm()` that computes the Euclidean norm (i.e. length) of a vector: $\|\mathbf{x}\| = \sqrt{\mathbf{x}^T \mathbf{x}}$.

```

# test vnorm() with 1:5
v <- 1:5
vnorm(v)

```

- Given the vector `v <- 1:5`, use `vnorm()` to create a vector `u` such that `u` is of unit norm: $\|\mathbf{u}\| = 1$.
- Write a function `is_square()` to check whether the provided matrix is a square matrix.
- Write a function `mtrace()` to compute the trace of a square matrix. Use your `is_square()` function to make sure that the input is a square matrix. The trace is defined as the sum of the elements on the diagonal: $tr(\mathbf{A}) = \sum_{i=1}^n a_{ii}$
- Given two square matrices `A` and `B`, verify that your function `tr()` is a linear mapping:
 - $tr(\mathbf{A} + \mathbf{B}) = tr(\mathbf{A}) + tr(\mathbf{B})$
 - $tr(c\mathbf{A}) = c \times tr(\mathbf{A})$, where c is a scalar

- Trace of products: Given two matrices `X` and `Y` of adequate size, verify that:

$$tr(\mathbf{X}^T \mathbf{Y}) = tr(\mathbf{X} \mathbf{Y}^T) = tr(\mathbf{Y}^T \mathbf{X}) = tr(\mathbf{Y} \mathbf{X}^T)$$

- Cross-products: Use `proc.time()` or `system.time()` to compare the `t(X) %*% X` against `crossprod(X)`
- Cross-products: Use `proc.time()` or `system.time()` to compare the `X %*% t(X)` against `tcrossprod(X)`

Transformation and Scaling Operations

In this section you will be using the built in data.frame `mtcars` to practice transforming and scaling operations:

```
head(mtcars)
```

```

##              mpg cyl  disp  hp  drat   wt  qsec vs am gear carb
## Mazda RX4      21.0   6  160 110  3.90 2.620 16.46  0  1    4    4
## Mazda RX4 Wag  21.0   6  160 110  3.90 2.875 17.02  0  1    4    4
## Datsun 710      22.8   4  108  93  3.85 2.320 18.61  1  1    4    1
## Hornet 4 Drive  21.4   6  258 110  3.08 3.215 19.44  1  0    3    1
## Hornet Sportabout 18.7   8  360 175  3.15 3.440 17.02  0  0    3    2
## Valiant         18.1   6  225 105  2.76 3.460 20.22  1  0    3    1

```

- Create a matrix `M` with variables `mpg`, `disp`, `hp`, `drat`, and `wt`.
- Use `apply()` to compute the vector containing the means of the columns in `M`
- Compute a matrix `Mc` of mean-centered data applying the function `scale()` on `M` (do NOT `scale = TRUE`).
- Confirm that variables in `Mc` are mean-centered by calculating the vector of column-means
- Use the function `sweep()` to mean-center `M` by *sweeping out* the vector of column means. Compare this result with `Mc` (you should get the same values).
- Compute a vector of column maxima.
- Use `sweep()` to scale the columns by dividing by the column maxima.
- Compute a matrix in which all columns are scaled to have minimum = 0 and maximum = 1
- Write a function `dummify()` that takes a factor or a character vector, and which returns a matrix with dummy indicators. Assuming that the factor has k categories (i.e. k levels), include an argument `all` that lets you specify whether to return k binary indicators, or $k - 1$ indicators. You should be able to call `dummify()` like this:

```
cyl <- factor(mtcars$cyl)
# all categories
CYL1 <- dummify(cyl, all = TRUE)
# minus one category
CYL2 <- dummify(cyl, all = FALSE)
```

- Write a function `crosstable()` that takes two factors, and which returns a cross-table between those factors. To create this function you should use your function `dummify()` to compute two dummy matrices, and then multiple them.

```
cyl <- factor(mtcars$cyl)
gear <- factor(mtcars$gear)
xtb <- crosstable(cyl, gear)
```

You should get a table like this:

```
      3 4 5
4    1 8 2
6    2 4 1
8   12 0 2
```

Solutions

```
M1 <- matrix(x, nrow = 3, ncol = 3)
```

```
M1
```

```
##      [,1] [,2] [,3]
## [1,]    1    4    7
## [2,]    2    5    8
## [3,]    3    6    9
```

```
M2 <- matrix(x, nrow = 3, ncol = 3, byrow = TRUE)
```

```
M2
```

```
##      [,1] [,2] [,3]
## [1,]    1    2    3
## [2,]    4    5    6
## [3,]    7    8    9
```

```
diag(1, nrow = 5)
```

```
##      [,1] [,2] [,3] [,4] [,5]
## [1,]    1    0    0    0    0
## [2,]    0    1    0    0    0
## [3,]    0    0    1    0    0
## [4,]    0    0    0    1    0
## [5,]    0    0    0    0    1
```

```
M3 <- cbind(a1, a2, a3)
```

```
rownames(M3) <- 1:nrow(M3)
```

```
M3
```

```
##   a1  a2 a3
## 1  2 1.88 80
## 2  3 2.05 90
## 3  6 1.70 70
## 4  7 1.60 50
## 5 10 1.78 75
```

```
M4 <- rbind(a1, a2, a3)
```

```
colnames(M4) <- 1:ncol(M4)
```

```
M4
```

```
##      1      2      3      4      5
## a1  2.00  3.00  6.0  7.0 10.00
## a2  1.88  2.05  1.7  1.6  1.78
## a3 80.00 90.00 70.0 50.0 75.00
```

```
# vector norm
```

```
vnorm <- function(x) {
  sqrt(t(x) %*% x)
}
```

```
vnorm(1:5)
```

```
##           [,1]
```

```
## [1,] 7.416198
```

```
# vector of unit norm
```

```
v <- 1:5
```

```
u <- v / vnorm(v)
```

```
vnorm(u)
```

```
##           [,1]
```

```
## [1,]      1
```

```
# is a matrix square
```

```
is_square <- function(m) {
```

```
  nrow(m) == ncol(m)
```

```
}
```

```
is_square(matrix(1:9, 3, 3))
```

```
## [1] TRUE
```

```
# function for trace of a matrix
```

```
mtrace <- function(m) {
```

```
  if (!is_square(m)) {
```

```
    stop("\n'mtrace()' requires a square matrix")
```

```
  }
```

```
  sum(diag(m))
```

```
}
```

```
mtrace(matrix(1:9, nrow = 3, ncol = 3))
```

```
## [1] 15
```

```
# trace is a linear mapping
```

```
set.seed(123)
```

```
A <- matrix(sample(1:9, 9), nrow = 3, ncol = 3)
```

```
B <- matrix(sample(1:9, 9), nrow = 3, ncol = 3)
```

```
mtrace(A + B)
```

```
## [1] 22
```

```
mtrace(A) + mtrace(B)
```

```
## [1] 22
```

```
mtrace(3 * A)
```

```
## [1] 36
```

```
3 * mtrace(A)
```

```
## [1] 36
# trace of a product
set.seed(123)
X <- matrix(sample(1:12, 12), nrow = 4, ncol = 3)
Y <- matrix(sample(1:12, 12), nrow = 4, ncol = 3)

mtrace(t(X) %*% Y)

## [1] 593
mtrace(X %*% t(Y))

## [1] 593
mtrace(t(Y) %*% X)

## [1] 593
mtrace(Y %*% t(X))

## [1] 593
# crossproduct
set.seed(345)
X <- matrix(runif(100000), 1000, 100)
system.time(t(X) %*% X)

##      user  system elapsed
##    0.007    0.000    0.007
system.time(crossprod(X))

##      user  system elapsed
##    0.005    0.000    0.004
# tcrossproduct
set.seed(345)
Y <- matrix(runif(100000), 100, 1000)
system.time(Y %*% t(Y))

##      user  system elapsed
##    0.007    0.001    0.008
system.time(tcrossprod(Y))

##      user  system elapsed
##    0.004    0.000    0.004
M <- as.matrix(mtcars[, c('mpg', 'disp', 'hp', 'drat', 'wt')])

# vector of column-means
means <- apply(M, MARGIN = 2, FUN = mean)

# mean-centered Mc with scale()
```

```

Mc <- scale(M, scale = FALSE)
apply(Mc, MARGIN = 2, FUN = mean)

##           mpg           disp           hp           drat           wt
## 4.440892e-16 -1.199041e-14  0.000000e+00 -1.526557e-16  3.469447e-17

# mean-centered with sweep()
Mc2 <- sweep(M, MARGIN = 2, STATS = means, FUN = "-")
apply(Mc2, MARGIN = 2, FUN = mean)

##           mpg           disp           hp           drat           wt
## 4.440892e-16 -1.199041e-14  0.000000e+00 -1.526557e-16  3.469447e-17

# scaling by the maximum
maxs <- apply(M, MARGIN = 2, FUN = max)
Mx <- sweep(M, MARGIN = 2, STATS = maxs, FUN = "/")

# rescaling from 0 to 1
mins <- apply(M, MARGIN = 2, FUN = min)
ranges <- apply(M, MARGIN = 2, FUN = function(x) max(x) - min(x))
M01 <- scale(M, center = mins, scale = ranges)
# test it
apply(M01, MARGIN = 2, FUN = min)

## mpg disp  hp drat  wt
##  0   0   0   0   0

apply(M01, MARGIN = 2, FUN = max)

## mpg disp  hp drat  wt
##  1   1   1   1   1

# function dummify
dummify <- function(x, all = FALSE) {
  if (!is.factor(x)) x <- as.factor(x)
  categories <- levels(x)
  num_categories <- length(categories)
  if (!all) {
    num_categories <- length(categories) - 1
  }
  dummies <- matrix(0, nrow = length(x), ncol = num_categories)
  for (k in 1:num_categories) {
    dummies[x == categories[k],k] <- 1
  }
  colnames(dummies) <- categories[1:num_categories]
  dummies
}

# test it
cyl <- factor(mtcars$cyl)

```

```
CYL1 <- dummify(cyl, all = TRUE)
CYL2 <- dummify(cyl, all = FALSE)

# function crosstable
crosstable <- function(x, y) {
  if (!is.factor(x)) x <- factor(x)
  if (!is.factor(y)) y <- factor(y)
  Xdum <- dummify(x, all = TRUE)
  Ydum <- dummify(y, all = TRUE)
  t(Xdum) %*% Ydum
}

cyl <- factor(mtcars$cyl)
gear <- factor(mtcars$gear)
xtb <- crosstable(cyl, gear)
```