

Stat 243

R Programming Concepts

Gaston Sanchez

Creative Commons Attribution 4.0 International License

Programming Concepts

References

Most of the content in these slides is a curated collection of notes and quotes from:

- ▶ **R Programming for Bioinformatics** by Robert Gentleman.
- ▶ **Software for Data Analysis** by John Chambers.
- ▶ **Extending R** by John Chambers.
- ▶ **Advanced R** by Hadley Wickham.
- ▶ **The Art of R Programming** by Norman Matloff.
- ▶ **Introduction to Scientific Programming and Simulation Using R** by O. Jones, R. Maillardet, and A. Robinson.

R sessions

A typical session using R involves:

- ▶ starting R
- ▶ loading packages necessary to perform intended analysis
- ▶ loading data sets
- ▶ data manipulation in a variety of ways
- ▶ perhaps saving results and output to an external location
- ▶ *etc*

R sessions

- ▶ For every R session you have a **workspace**.
- ▶ The workspace is referred to as the **global environment**.
- ▶ This is the place where the objects that you create will be stored.
- ▶ The function `ls()` lists the objects in your current workspace.
- ▶ More information about your session is provided with `sessionInfo()`.

R language basics

- ▶ Programming in R is carried out by manipulating and modifying data structures.
- ▶ These different transformations are carried out using functions and operators.
- ▶ In R, virtually every operation is a function call.
- ▶ In many ways, R can be considered a functional programming language.
- ▶ Among other things, functions can be passed as arguments to functions, and returned as values.

Pretty much everything is a function call

Consider the following code

```
a <- 1:5  
b <- 2  
a + b
```

- ▶ The assignment operator "<-" is a function
- ▶ The colon operator ":" is also a function
- ▶ The "+" operator is another function

Pretty much everything is a function call

```
# colon operator as a function
```

```
":"(from = 1, to = 5)
```

```
# assignment as a function
```

```
"<-"(a, value = ":"(from = 1, to = 5))
```

```
# addition as a function
```

```
"+"(2, 3)
```


Objects and Names

- ▶ As in any programming language, it's essential to be able to refer to objects, in a particular context, in a way that is consistent and clear.
- ▶ In R, there is one and only one way to refer to objects: by **name and an environment** (i.e. context)

Objects and Names

```
x <- 1
```

- ▶ Say you create a variable (i.e. object) `x`
- ▶ You use the assignment operator `<-` (or `=`) for this purpose
- ▶ Formally, this object has the name "x"
- ▶ The value of 1 is **assigned** to the object `x`
- ▶ `x` lives in the **global environment** (i.e. your workspace)
- ▶ `x` will be given a memory address

Objects and references

To see the address of `x`, use `tracemem()`

```
# memory address  
tracemem(x)
```

```
## [1] "<0x7ff27cabff78>"
```

Objects and references

Unless you make an assignment to `x`, it will keep the same address:

```
x + 2
```

```
## [1] 3
```

```
log(x)
```

```
## [1] 0
```

```
tracemem(x) # same memory address as before
```

```
## [1] "<0x7ff27cabff78>"
```

Objects and references

When you make an assignment to `x`, it will change address:

```
# assigning a name to x (will modify its memory address)  
names(x) <- "a"
```

```
## tracemem[0x7ff27cabff78 -> 0x7ff27aa7a0f8]: eval eval w
```

Objects and Environments

- ▶ Any computation in R takes place in an environment.
- ▶ An environment defines how the evaluator will search for an object by name.
- ▶ The part of a program in which an object is defined is called its **scope**.
- ▶ We talked about environments when discussing function evaluation.

Assignments and Replacements

Assignments

As you know, you can use the arrow operator "<-" to create objects, and also to modify them:

```
x <- c(1, 2, 3)
x
```

```
## [1] 1 2 3
```

```
x[1] <- 0
x
```

```
## [1] 0 2 3
```


Assignments

an assignment

```
x <- c(1, 2, 3)
```

is actually a call to the assignment function "<-"()

```
"<-"(x, value = c(1, 2, 3))
```

Replacement

a replacement

```
x[1] <- 0
```

is actually a call to the replacement function "[<-"()

```
x <- "[<-"(x, 1, value = 0)
```

What happens in a replacement?

```
# replacement  
x[1] <- 0
```

- ▶ `x[1] <- 0` gives the impression of having changed the value of `x`.
- ▶ For practical purposes, you can say that “we changed the value of `x`”.
- ▶ However, R didn't change `x` directly.
- ▶ Instead, a copy of `x` was created first, and then the modification took place in such copy.
- ▶ After that, the name “`x`” was assigned to the modified copy.

What happens in a replacement?

In R, you constantly use expressions like this:

```
z <- runif(4)
z
```

```
## [1] 0.1589814 0.4758246 0.7036687 0.8047777
```

```
z <- z - mean(z)
z
```

```
## [1] -0.37683165 -0.05998855 0.16785561 0.26896459
```

What happens in a replacement?

Consider the replacement expression: `z <- z - mean(z)`

- ▶ a replacement creates a new assignment of an object to the current name
- ▶ a new complete object `z` replaces the existing `z`
- ▶ this means that R makes a copy of `z`

Replacements and Environments

- ▶ An object assigned in an environment would only be changed by another assignment:

```
y <- 1:3  
tracemem(y)
```

```
## [1] "<0x7ff27cd210e8>"
```

```
y[3] <- 5
```

```
## tracemem[0x7ff27cd210e8 -> 0x7ff27ccf3278]: eval eval w  
## tracemem[0x7ff27ccf3278 -> 0x7ff27ccdde50]: eval eval w
```

Copy-on-change policy

x and a copy of x, share the same address

```
x <- c(1, 2, 3)
```

```
y <- x
```

```
tracemem(x)
```

```
## [1] "<0x7ff27ca7aa38>"
```

```
tracemem(y)
```

```
## [1] "<0x7ff27ca7aa38>"
```

What if you modify x?

```
x[3] <- 4
```

```
## tracemem[0x7ff27ca7aa38 -> 0x7ff27c704728]: eval eval w
```

```
tracemem(x)
```

```
## [1] "<0x7ff27c704728>"
```

```
tracemem(y)
```

```
## [1] "<0x7ff27ca7aa38>"
```

x had an initial address, then was copied to another address, and ultimately it ended in <0x7ff27c704728>

Copy-on-change policy

```
# what if you modify x?  
x[3] <- 4
```

Even though we are changing just one element of the vector `x`, the entire vector is reassigned. This means that the entire vector is copied, which can take up a lot of time for long vectors (especially inside a loop).

Boxing and Unboxing

how does R handle this addition?

```
x <- c(2, 4, 6)
```

```
y <- c(3, 5, 7)
```

```
z <- x[1] + y[1]
```

Boxing and Unboxing




Figure 1: illustration of boxing and unboxing

Boxing and Unboxing

- ▶ Every value in R must be stored in a vector.
- ▶ This is true for values which look like scalars such as `TRUE`, `17`, and `"hello"`.
- ▶ Operations which are specified element-by-element on vectors occur overhead because the elements must be extracted into a vector of length 1, processed, and then transferred back to a longer vector.

Memory Management

Memory allocation

- ▶ Memory management in R is automatic
- ▶ Users do not explicitly manage memory for R objects: you do not need to allocate or delete storage explicitly
- ▶ All objects are dynamically allocated at some time during the R session
- ▶ Evaluating function calls and assigning objects triggers allocation and occasional deallocation (via garbage collection)

Memory and garbage collection

- ▶ In particular, objects cannot be deleted
- ▶ You can use the function `rm()` to remove the assignment of an object in an environment, but does not delete the memory of the object
- ▶ Unused dynamic storage is recovered by *garbage collection*
- ▶ Garbage collection takes place automatically and unpredictably

Memory and garbage collection

- ▶ Garbage collector runs automatically whenever storage requests exceed the current free memory available
- ▶ It is possible to force garbage collection with the function `gc()`
- ▶ You can use `gcinfo()` to print information when the garbage collector runs
- ▶ Garbage collection recovers storage from objects that do not have a current reference

R's Memory Management

- ▶ In R, memory is divided into 2 separate components:
 - memory for atomic vectors
 - memory for other language elements

R's Memory Management

- ▶ In R, memory is divided into 2 separate components:
 - memory for atomic vectors
 - memory for other language elements
- ▶ Vector storage is contiguous storage for homogenous elements (i.e. atomic)
- ▶ Vector storage is further divided into two types:
 - small vectors (less than 128 bytes) which are allocated in R
 - large vectors for which memory is obtained from the OS

Efficient Code

R language basics

To use R effectively, it is important to understand which operations are likely to be “expensive” and how the expense can be avoided.

The main expensive operations are:

- ▶ Boxing and unboxing
- ▶ Data copying
- ▶ Argument matching (when calling a function)

Data Copying

- ▶ When a value is passed into a function, R works hard to make it appear that it is a copy of the value which has been handed to the function.
- ▶ Many of the copy operations which R performs are unnecessary because code which decides whether it is necessary to copy errs on the conservative side.

Argument Matching

- ▶ The presence of optional arguments to R functions means that function calls are relatively expensive. (Three passes are made over the argument list).
- ▶ Some functions (mainly the simple mathematical functions) avoid the overhead of argument processing.
- ▶ These functions are called *primitive* and can be recognised by their appearance.

R language basics

- ▶ The R evaluator is written in C.
- ▶ Many commands are written in C and thus do run in machine code.
- ▶ But most other functions (and the ones you write) are written in R itself.
- ▶ Those commands that use R code are thus interpreted.
 - R is an interpreted language.

Functions written in C

Functions written in C are indicated with `.Primitive()`

```
# addition operator  
get("+")
```

```
## function (e1, e2) .Primitive("+")
```

```
# square root  
sqrt
```

```
## function (x) .Primitive("sqrt")
```


Writing “optimized” R code

- ▶ Whenever possible use **Vectorization**
- ▶ Understand R's functional programming nature
- ▶ Understand how R uses memory
- ▶ Use some of R's optimized utility functions

Looking at for loops

```
x <- 1:10
y <- 21:30

# for-loop
z <- NULL
for (i in 1:length(x)) {
  # at each iteration we "expand" z with c()
  # (z gets copied and increases its memory)
  z <- c(z, x[i] + y[i])
}
```

Looking at for loops

Adding two vectors element by element

```
x <- 1:10
y <- 21:30

# for-loop
z <- NULL
for (i in 1:length(x)) {
  # at each iteration we "append" a new element to z
  # (z gets copied and increases its memory)
  z[i] <- x[i] + y[i]
}
```

Looking at for loops

```
x <- 1:10
y <- 21:30

# for-loop
z <- numeric(0)
for (i in 1:length(x)) {
  # at each iteration we "append" a new element to z
  # (z gets copied and increases its memory)
  z[i] <- x[i] + y[i]
}
```

Looking at for loops

```
x <- 1:10
y <- 21:30

# better to "allocate" 'z' with appropriate length
z <- numeric(10)
for (i in 1:length(x)) {
  # at each iteration we "append" a new element to z
  # (z gets copied without increasing its memory)
  z[i] <- x[i] + y[i]
}
```

R for loops

```
z <- numeric(0)
for (i in 1:length(x)) {
  z[i] <- x[i] + y[i]
}
```

- ▶ `for()` is a function
- ▶ the colon operator `:` is also a function
- ▶ vector subscript operation `[]` is also a function
- ▶ vector subscript assignment `[]=` is also a function
- ▶ each iteration allocates a new object `z`

R for loops

```
x <- 1:10
y <- 21:30

# there are many function calls inside this loop
z <- numeric(10)
for (i in 1:length(x)) {
  z[i] <- x[i] + y[i]
}
```

About for loops

- ▶ R for loops have a bad reputation of being “slow”
- ▶ It is not for loops that are inherently slow
- ▶ It is what you do inside a loop that may be slow
- ▶ Especially when you perform many replacement operations

```
x <- 1:10  
y <- 21:30
```

```
# whenever possible use vectorized code
```

```
z <- x + y
```


Vectorized functions

Many functions in R, and most functions that work with vectors, are already vectorized:

- ▶ arithmetic
- ▶ trigonometric
- ▶ set operations
- ▶ comparison operators
- ▶ logical operators

Vectorized Functions

Other vectorized functions are:

- ▶ `ifelse()`
- ▶ `which()`
- ▶ `where()`
- ▶ `any()`
- ▶ `all()`

Vectorized matrix operations

- ▶ `rowSums()`
- ▶ `colSums()`
- ▶ `rowMeans()`
- ▶ `colMeans()`

PS1 Example

A room contains 100 toggle switches, originally all turned off:

```
num_switches <- 100  
switches <- rep("off", num_switches)
```

- ▶ 100 people enter the room in turn.
- ▶ the first person toggles every switch
- ▶ the second toggles every second switch
- ▶ the third every third switch, and so on,
- ▶ to the last person who toggles the last switch only.
- ▶ Which switches are turned on?

PS1 Example: most common answer

```
num_switches <- 100
switches <- rep("off", num_switches)

for (i in 1:num_switches) {
  positions <- seq(i, num_switches, by = i)
  for (pos in positions) {
    if (switches[pos] == "off") {
      switches[pos] <- "on"
    } else {
      switches[pos] <- "off"
    }
  }
}
```

PS1 Example: with vectorized ifelse()

```
switches <- rep("off", num_switches)

for (i in 1:num_switches) {
  positions <- seq(i, num_switches, by = i)
  switches[positions] <- ifelse(
    switches[positions] == "off", "on", "off")
}
```

PS1 Example: my solution with a “lookup” vector

```
switches <- rep("off", num_switches)

# ProTip: handy "lookup" vector
toggle <- c("on" = "off", "off" = "on")

for (i in 1:num_switches) {
  positions <- seq(i, num_switches, by = i)
  switches[positions] <- toggle[switches[positions]]
}
```

Tools to time R code

There are various handy functions that can help you evaluate and profile your code:

- ▶ `proc.time()`: determines how much real and CPU time (in seconds) the currently running R process has already taken.
- ▶ `system.time()`: measures CPU time used in an **R expression**
- ▶ `Rprof()`: profiling of the execution of R expressions
- ▶ `microbenchmark()`: from package "microbenchmark"

Function `proc.time()`

- ▶ `proc.time()` works as a stop-watch
- ▶ you initialize it to a starting time and then run all the desired code
- ▶ you stop it by subtracting the starting time from the ending time

CPU time `proc.time()`

```
g <- rnorm(100000)
h <- rep(NA, 100000)

# Start the clock
proc_time <- proc.time()

# Loop through the vector, adding one
for (i in 1:100000){
  h[i] <- g[i] + 1
}

# Stop the clock
proc.time() - proc_time
```

CPU time `proc.time()`

```
g <- rnorm(100000)
h <- rep(NA, 100000)

# Start the clock
proc_time <- proc.time()
# vectorized addition
h <- g + 1
# Stop the clock
proc.time() - proc_time
```

```
##      user  system elapsed
##    0.002    0.000    0.002
```

How to read `proc.time()` output?

- ▶ The values `user`, `system`, `elapsed` will be defined by your operating system.
- ▶ The `user` time relates to the execution of the code.
- ▶ The `system` time relates to your CPU, and it is the time spent in system calls.
- ▶ The `elapsed` time is the difference in times since you started the stopwatch. This is the actual time it took for the expression to be evaluated.
- ▶ It is usually the elapsed time which is important.

Function `system.time()`

- ▶ `system.time()` is similar to `proc.time()`
- ▶ However, `system.time()` takes a single R expression as its argument
- ▶ Typically you encapsulate your code in a function and pass it to `system.time()`

Timing code with `system.time()`

```
fast_add <- function(x) {  
  return(x + 1)  
}
```

```
system.time(fast_add(g))
```

```
##      user  system elapsed  
##    0.000    0.001    0.000
```

Timing code with `system.time()`

```
loop_add <- function(x) {  
  y <- rep(NA, length(x))  
  for (i in 1:length(x)){  
    y[i] <- x[i] + 1  
  }  
  return(h)  
}  
  
system.time(loop_add(g))
```

```
##      user  system elapsed  
##    0.100    0.004    0.108
```

Monitoring with Rprof()

```
num_switches <- 10000  
switches <- rep("off", num_switches)
```

```
Rprof()
```

```
for (i in 1:num_switches) {  
  positions <- seq(i, num_switches, by = i)  
  for (pos in positions) {  
    if (switches[pos] == "off") {  
      switches[pos] <- "on"  
    } else {  
      switches[pos] <- "off"  
    }  
  }  
}
```

```
Rprof(NULL)
```

```
summary_rprof <- summaryRprof()
```


How `Rprof()` works?

- ▶ R inspects the call stack to determine which function calls are in effect at that time
- ▶ It writes the result of each inspection to a file: `Rprof.out`
- ▶ The function `summaryRprof()` conveniently summarizes all those lines in the file
- ▶ If the code you're profiling produces many function calls, the profiling output may be hard to decipher.

Some examples from Matloff (2011)

Expected value of the maximum of independent $N(0,1)$ random variables X and Y :

```
sum <- 0
num_reps <- 10000
for (i in 1:num_reps) {
  xy <- rnorm(2)  # generate 2 N(0,1)s
  sum <- sum + max(xy)
}
sum / num_reps
```

```
## [1] 0.5785969
```

Some examples from Matloff (2011)

Expected value of the maximum of independent $N(0, 1)$ random variables X and Y :

- ▶ generated 10,000 pairs with: `xy <- rnorm(2)`
- ▶ use a for loop to find the maximum for each pair: `max(xy)`
- ▶ compute estimated expected value: `sum / num_reps`

Some examples from Matloff (2011)

```
sum <- 0

num_reps <- 10000

for (i in 1:num_reps) {
  xy <- rnorm(2)  # generate 2  $N(0,1)$ s
  sum <- sum + max(xy)
}
```

- ▶ at each iteration we call `rnorm()` and `max()`
- ▶ at each iteration there is an assignment operation

A faster implementation

```
num_reps <- 10000

xy_matrix <- matrix(rnorm(2 * num_reps), ncol = 2)

maximums <- pmax(xy_matrix[,1], xy_matrix[,2])

mean(maximums)

## [1] 0.5719855
```

Some Programming Habits

- ▶ Good programming is clear rather than clever
- ▶ Being clever is good, but given the choice, being clear is preferable
- ▶ You can start addressing a problem with a `for` loop that is *clear*
- ▶ Once you know the code works, then see what operations can be vectorized
- ▶ Identifying vectorized options takes time and practice

Lessons

- ▶ Optimizing R performance can be a very dark art.
- ▶ Knowing about the detail of how functions work internally can be helpful but is not essential.
- ▶ Experimentation with the code and timing the results with `system.time()` can reduce times by orders of magnitude.
- ▶ In general, vectorization is a big win.
- ▶ Converting loops into vectorized alternatives almost always pays off.
- ▶ Code profiling can give you a way to locate those parts of a program which will benefit most from optimization.